


# **A Chained Model for Weather Forecasting and Wildfire Risk Prediction**

**Team 2**

2491015 Bang Gaeun  
2491017 Baek Gyurie  
2491025 Lee Jiwon  
2491029 Jeong Ahhyeon

# TABLE OF CONTENTS



- |    |   |    |   |
|----|---|----|---|
| 01 | <b>Introduction</b>                             |    |   |
| 02 | <b>Model 1:<br/>preprocessing</b>               | 03 | <b>Model 1:<br/>Modeling &amp; Evaluation</b> |
| 04 | <b>Model 2:<br/>preprocessing</b>               | 05 | <b>Model 2:<br/>Modeling &amp; Evaluation</b> |
| 06 | <b>Deployment:<br/>External Function Design</b> | 07 | <b>Significance of the project</b>            |

# Introduction



It is necessary to predict the possibility of forest fires and prepare systematically!

**Predict weather in Gyeongbuk**



**Calculate the probability of  
forest fire occurrence**


02

## **Model 1:** preprocessing


## 02 Model 1: preprocessing

### Gyeongbuk monthly weather dataset(1949 - 2025)


6988 rows X 6 columns



- **Average maximum temperature**
- **Average minimum temperature**
- **Monthly precipitation**
- **Average wind speed**



#### Outliers(IQR)



5%

1%

#### Missing Values

0.24%

Too insignificant



Delete the rows

Rows: 6988 → 6971

If outliers are removed from the weather dataset, the model may not predict heavy rain, cold, and heat waves.



Remain the rows



## 02 Model 1: preprocessing

### Original dataset

지점	지점명	일시	평균최고기온(°C)	평균최저기온(°C)	월합강수량(00~24h)mm	평균풍속(m/s)
115	울릉도	1938-08			105.2	
115	울릉도	1938-09	22.7	17.9	153.1	4.5
115	울릉도	1938-10	18.2	13.6	81.6	5.6
115	울릉도	1938-11	11.6	6.4	196.1	5.7
115	울릉도	1938-12	5.6	0.7	70.3	4.6
115	울릉도	1939-01	2.8	-2.5	87.5	4.7
.						
283	경주시	2025-03	15.3	2.4	55.7	3.2
283	경주시	2025-04	21.6	6.5	34.4	3.1

### Remove Ulleungdo data

Based on domain knowledge, Ulleungdo has a remarkably different climate alone.

### Divide the date column by year and month

To predict the monthly weather

### Month → Categorical Encoding

### Unify features in Dataset 1 with Dataset 2

### Unify feature names with dataset 2

Column name Korean → English

### Generate 'temp\_range' variable

Unify feature with dataset 2

```
df['temp_range'] = df['MAX_TEMP'] - df['MIN_TEMP']
```

### Regional Aggregation for Province-Level Climate Representation

To generate region-wide representative climate data for the Gyeongbuk region, we averaged the values across all constituent cities and counties monthly and yearly. After aggregation, we removed the original location information, which was no longer necessary.

### Monthly precipitation (월합강수량)

→ conversion to daily average precipitation



Since Dataset 2 has an daily precipitation feature, we conducted a process to match it. Monthly precipitation is the addition of all the daily precipitation each month. We divided it by the number of days each month and converted it into average daily precipitation.

## 02 Model 1: preprocessing

### Tasks to improve the accuracy of weather prediction

#### 01 Use as learning data since 1955

When plotting each feature by year, the data from 1949 to 1954 tended to show a remarkably different distribution. Therefore, deleting the data during this period could improve the model's performance.



#### 02 Replace 'YEAR' column with 'YEAR\_SINCE\_2000'

If we use the year as it is, the difference between the years is too small compared to the year's number size, which is a thousand digits. Therefore, we designated 2000 as the base year and added the 'difference from the base year' as the new column.



## 02 Model 1: preprocessing

### Tasks to improve the accuracy of weather prediction

#### 03 Create three features through Feature Engineering

These three derivative features allow the model to learn well monthly and year-to-year differences in weather.

##### 1. 'LINEAR\_TEMP\_TREND'

Contains information that global warming is gradually raising temperatures.

```
linreg = LinearRegression()  
linreg.fit(df[['YEAR_SINCE_2000']], df['MAX_TEMP'])  
print(linreg.coef_) # 연도당 기온 변화량(=기울기)
```

[0.01714123]

df['LINEAR\_TEMP\_TREND'] = 0.017 \* df['YEAR\_SINCE\_2000']

We calculated the trend of temperature change per year in the Gyeongbuk dataset and provided it as domain knowledge.

##### 2. 'YEAR\_MONTH\_INDEX'

Expresses the monthly time flow in integers. This feature allows the model to learn temporal continuity.

```
df['YEAR_MONTH_INDEX'] = df['YEAR_SINCE_2000'] * 12 + df['MONTH']
```

##### 3. 'YEAR\_SQUARED'

stands for nonlinear temporal characteristics. This feature reflects the assumption that climate change may not be linear over the years.

```
df['YEAR_SQUARED'] = df['YEAR_SINCE_2000'] ** 2
```

## 02 Model 1: preprocessing

### Final 21 Features

MAX\_TEMP  
MIN\_TEMP  
AVG\_WIND\_SPEED  
temp\_range  
daily\_precip  
MONTH\_1  
MONTH\_2  
MONTH\_3  
  
.  
  
.  
  
MONTH\_12  
YEAR\_SINCE\_2000  
LINEAR\_TEMP\_TREND  
YEAR\_MONTH\_INDEX  
YEAR\_SQUARED

### Split the Dataset - Holdout Validation


#### Holdout Validation

verifies future data  
after a specific point in time.

#### K-fold Validation

divides data at various points in time  
and performs cross-validation.

Since Model 1 has to predict the future based on past data,  
the holdout method is more suitable than the k-fold method.



03

# **Model 1: Modeling & Evaluation**

# Overview of the Model 1



In Model 1, input features were designed to account for both **seasonality** and **long-term trends** in order to accurately predict climate conditions.

## Input

YEAR\_SINCE\_2000, YEAR\_SQUARED,  
YEAR\_MONTH\_INDEX, MONTH\_1 to MONTH\_12,  
LINEAR\_TEMP\_TREND

## Output

MAX\_TEMP, MIN\_TEMP, AVG\_WIND\_SPEED, temp\_range,  
and daily\_precip(total of 5 variables)



# CatBoostRegressor

## - Ordered Boosting

- : Better captures time-based patterns(ex. Year, Month)
- Shows greater stability than other GBDT models
  - : More reliable for long-term forecasts (e.g., 2030, 2080)
- Naturally handles **categorical features**
- Performs well even with a **small number of features**
- Enables fast experimentation without heavy hyperparameter tuning



## Extrapolation Weakness & Multi-Target Prediction Structure

To compensate for CatBoost's extrapolation weakness, a stacking structure combining CatBoost with extrapolation-robust Linear Regression

## MultiOutputRegressor

1. Creates and trains separate models(CatBoost+Linear+RidgeCV) for each target and then combines their predictions
2. Multiple target variables can be saved and used as a single unified model.

# Model 1 Tuning



## CatBoost

### Optuna

A single set of hyperparameters was optimized (for all target variables)

Optimized set was then replicated across individual CatBoost models for each target variable.



## StackingRegressor

### Base models:

[Shared CatBoost,  
Target-specific LinearRegression]

Meta-model: RidgeCV  
(automatically adjusts the regularization parameter to prevent overfitting)



## Linear Regression

Trained for each target variable (e.g., one for MAX\_TEMP, another for MIN\_TEMP, etc.).



## MultiOutputRegressor

All per-target StackingRegressor models were wrapped with a MultiOutputRegressor

# Model 1

- Optuna Hyperparameter Search

```
# Optuna objective function
def objective(trial):
    params = {
        'iterations': trial.suggest_int('iterations', 500, 2000),
        'depth': trial.suggest_int('depth', 4, 10),
        'learning_rate': trial.suggest_loguniform('learning_rate', 0.01, 0.3),
        'l2_leaf_reg': trial.suggest_float('l2_leaf_reg', 1, 10),
        'random_seed': 42,
        'verbose': 0,
    }
```

- Training and Validation of Stacking Models

```
# Train and predict for each target variable
val_preds = []
for target in target_cols:
    # CatBoost model
    cb = CatBoostRegressor(**params)

    # Base models: common CatBoost + individual LinearRegression
    base_models = [
        ('catboost', cb),
        ('linear', LinearRegression())
    ]
    meta_model = RidgeCV()

    stack = StackingRegressor(estimators=base_models, final_estimator=meta_model)
    stack.fit(X_train, y_train[target])
    val_pred = stack.predict(X_val)
    val_preds.append(val_pred)
```

- Performance Evaluation and Best Parameters Selection

```
# Calculate mean MAE across all targets
val_preds = np.array(val_preds).T # (samples, targets)
mae = np.mean([mean_absolute_error(y_val[col], val_preds[:, i]) for i, col in enumerate(target_cols)])
return mae
```

```
# Parameter tuning with Optuna
study = optuna.create_study(direction='minimize')
study.optimize(objective, n_trials=100)
```

```
# Fix best parameters
best_params = study.best_params
best_params.update({'random_seed': 42, 'verbose': 0})
```

## Optuna Hyperparameter Tuning

## Stacking model

# Model 1

- Stacking Models Trained on Full train set

```
# Full dataset by combining training and validation sets
X_full_train = pd.concat([X_train, X_val], axis=0)
y_full_train = pd.concat([y_train, y_val], axis=0)

# Predict on the test set
test_preds = []

cb = CatBoostRegressor(**best_params)

# Define stacking model
base_models = [
    ('catboost', cb),
    ('linear', LinearRegression())
]
meta_model = RidgeCV()
stacking_model = StackingRegressor(estimators=base_models, final_estimator=meta_model)

multioutput_model = MultiOutputRegressor(stacking_model)
multioutput_model.fit(X_full_train, y_full_train)

test_preds = multioutput_model.predict(X_test) # Shape: (n_samples, n_targets)
test_results = evaluate_all_targets(y_test, test_preds, target_cols, dataset_name="Test")
```

## Stacking model

# Model1: Evaluation

## Temperature-related variables (MAX\_TEMP, MIN\_TEMP)

→ The model predicts these variables very well(effective for temperature-related data)

## AVG\_WIND\_SPEED, daily\_precip:

→ Harder to predict due to their irregular and external factors, resulting in lower accuracy.

## Performance of the Stacking Model:





→ Has limitations when dealing with more discontinuous variables or those heavily influenced by external factors, such as precipitation

→ Shows strong performance on continuous variables with clear seasonal patterns, such as temperature

Target	MAE	MSE	RMSE	1 - R <sup>2</sup>
MAX_TEMP	1.1548	2.2610	1.5036	<b>0.0299</b>
MIN_TEMP	1.0940	2.0284	1.4242	<b>0.0224</b>
AVG_WIND_SPEED	0.1914	0.0590	0.2428	0.4260
temp_range	0.8424	1.0727	1.0357	0.2964
daily_precip	1.5228	4.1469	2.0364	0.4063

# Visualization

## Temperature-related variables



04

## **Model 2: preprocessing**

# Removing Irrelevant Features

- The second dataset was sourced from Zenodo and provides a comprehensive compilation of weather observations and wildfire occurrences in California from 1984 to 2025


Removed variables such as date, cumulative precipitation, and others not aligned with Model 1's outputs.

Only kept

- Meteorological features that overlap with Model 1's outputs. (e.g., MAX\_TEMP, AVG\_WIND\_SPEED)
- Season (instead of month)
- Wildfire occurrence

Month was excluded due to climate differences between California and Korea.

Instead, we used season categories to reflect general temporal patterns



# Data Cleaning & Scaling



## Missing Values

- The dataset contains 14,988 rows in total.
- Among features, only 'AVG\_WIND\_SPEED' had 12 missing values. (The missing rate is extremely low)
- Therefore, **rows with missing values were simply removed.**



## Outlier (IQR)

- Variables except 'PRECIPITATION' had outlier rates under 4%.
- 'PRECIPITATION' had 9.22% outliers. However, this is due to many values being zero, which is natural for this variable. (The 75th percentile is also 0mm.)
- Random Forest is robust to outliers
- **No outlier removal was applied.**



## No Scaling

- Random Forest models are not sensitive to feature scaling.
- Therefore, **no normalization or standardization was applied.**

# Data Transformation

**Unit Conversion:** To match Dataset 1, all units were standardized.

Feature	From	To
Temperature	°F	°C
Precipitation	inches	mm
Wind Speed	mph	m/s


**Encoding:** Season was also encoded as a categorical variable.

05

# Model 2: Modeling & Evaluation

# Overview of the Model 2

- **Input**=(Max temp, Min temp, Temp Range, Avg Wind Speed, Precipitation, Season)
- This is a **probabilistic prediction model**.
- We used **Random Forest Classifier**.
- Capable of capturing complex nonlinear relationships, Strong performance on tabular data.



# Data splitting and Model tuning



## Train Test Split

Dataset was split into

Training Set: 80%

Test Set: 20% (for evaluation)



## Train Validation Split

Training Set was split again into

Training Subset: 80%

Hold-out Validation Set: 20% (for comparing the performance of different tuning strategies and selecting the best-performing hyperparameters)



## Model tuning

Applied 3 different tuning strategies, including:

1. Random Search
2. Random Search + Grid Search
3. Bayesian Optimization



## Evaluation

Evaluating final model performance on test set.

# Model Tuning

- **Model tuning was conducted using 5-fold cross-validation within the training subset**, applying three different tuning strategies to identify the most effective approach for selecting optimal hyperparameters.
- **Compared the performances of the model using three different tuning strategies on a hold-out validation set** and selected the final model with the optimal hyperparameters.
- A custom scoring function was used, defined **as  $0.5 \times \text{ROC AUC} + 0.5 \times \text{PR AUC}$** , since the model is a probabilistic predictor rather than a binary classifier, and thus performance should not depend on a specific threshold.

# Model Tuning

- Random Search

```
param_dist = {  
    'n_estimators': [120, 150, 180, 200, 220],  
    'max_depth': [8, 10, 15, 20],  
    'min_samples_split': [8, 10, 12, 15],  
    'min_samples_leaf': [3, 4, 5],  
    'max_features': ['sqrt', 'log2']  
}  
}  
  
# Random Search  
random_search = RandomizedSearchCV(  
    estimator=rf,  
    param_distributions=param_dist,  
    n_iter=90,  
    scoring=custom_scorer,  
    cv=cv,  
    verbose=1,  
    random_state=42,  
    n_jobs=-1,  
    error_score='raise'  
)  
o  
o
```

- Random Search + Grid Search

```
# Parameter range refinement based on random search results  
param_grid = {  
    'n_estimators': [120, 150, 180, 200, 230],  
    'max_depth': [10, 12, 15, 18],  
    'min_samples_split': [12, 15, 18, 20],  
    'min_samples_leaf': [4, 5, 6],  
    'max_features': ['sqrt', 'log2']  
}  
}
```

```
# Grid Search  
grid_search = GridSearchCV(  
    estimator=rf,  
    param_grid=param_grid,  
    scoring=custom_scorer,  
    cv=cv,  
    n_jobs=-1,  
    verbose=1,  
    error_score='raise'  
)
```

- Bayesian Optimization

```
param_space = {  
    'n_estimators': Integer(100, 300),  
    'max_depth': Integer(5, 50),  
    'min_samples_split': Integer(2, 20),  
    'min_samples_leaf': Integer(1, 10),  
    'max_features': Categorical(['sqrt', 'log2']),  
    'bootstrap': Categorical([True, False])  
}  
}
```

```
bayes_search = BayesSearchCV(  
    estimator=rf,  
    search_spaces=param_space,  
    scoring=custom_scorer,  
    cv=cv,  
    n_iter=150,  
    n_jobs=-1,  
    verbose=1,  
    random_state=42  
)
```

# Model Tuning

- Random Search

- Best parameters:

```
'n_estimators': 120,  
'min_samples_split': 15,  
'min_samples_leaf': 5,  
'max_features': 'sqrt',  
'max_depth': 10
```

- Random Search + Grid Search


- Best parameters:

```
'n_estimators': 180,  
'min_samples_split': 18,  
'min_samples_leaf': 5,  
'max_features': 'sqrt',  
'max_depth': 12
```

- Bayesian Optimization

- Best parameters:


```
'n_estimators': 264,  
'min_samples_split': 2,  
'min_samples_leaf': 10,  
'max_features': 'log2',  
'max_depth': 50
```



# Model Tuning

	Random Search	Random + Grid	Bayesian
<b>CV Score</b>	0.6361	<b>0.6378</b>	0.6348
<b>PR AUC</b> (on hold-out validation set)	<b>0.6724</b>	0.6705	0.6677
<b>ROC AUC</b> (on hold-out validation set)	<b>0.8285</b>	0.8265	0.8273
<b>0.5 PR AUC + 0.5 ROC AUC</b> (on hold-out validation set)	<b>0.7505</b>	0.7485	0.7475

Based on the performance comparison, the optimal parameters obtained through **random search tuning were selected as the final hyperparameters**. Using these, the final Random Forest model was defined and **trained on the entire training set** (training subset + hold-out validation set). The trained model was then **evaluated on the test set**.



# Model 2 Performance

## Supplementary Metrics (Threshold = 0.5)

**Accuracy**  
**: 0.76502**

Measures the overall correctness of the model's wildfire predictions.

**Precision**  
**: 0.675958**


Indicates how many predicted wildfires were actually true.

**F1 Score**  
**: 0.623126**

Balances precision and recall to assess model performance.

**Recall:**  
**0.577954**

Shows how well the model catches actual wildfire occurrences. The model successfully identified 76% of actual wildfire occurrences.



# Model 2 Performance

## Key Evaluation Metrics

**Brier Score**  
**: 0.160779**


1. The closer to 0, the better
2. It is well-suited for evaluating the accuracy of probabilistic predictions

**Log Loss**  
**: 0.485913**

1. It evaluates how close the predicted probability is to the actual outcome.
2. Predictions that are close to 1(or 0) for the correct class receive better scores.


# **ROC AUC:** **0.819453**

Reflects the model's ability to distinguish wildfires from non-wildfires.



# **PR AUC:** **0.659028**


Evaluates performance in detecting rare events like wildfires.



# **Deployment**

## **06 : External Function Design**

# EXTERNAL FUNCTION



# GET INPUT

```
1 # Get input from the user
2 def get_user_input():
3     year = int(input("예측할 연도를 입력하세요 (예: 2026): "))
4     month = int(input("예측할 월을 입력하세요 (1~12): "))
5     return year, month
```

## Creating derived features & One-hot encoding

```
1 # One-hot encoding for Model 1
2 # Transform input data for Model 1
3 # Apply one-hot encoding to the MONTH column and adjust the order and missing columns to match model1_columns
4 def convert_to_model1_input(year, month, model1_columns):
5     df = pd.DataFrame([[year, month]], columns=['YEAR', 'MONTH'])
6
7     df['YEAR_SINCE_2000'] = year - 2000
8     df['LINEAR_TEMP_TREND'] = 0.017 * df['YEAR_SINCE_2000']
9     df['YEAR_MONTH_INDEX'] = df['YEAR_SINCE_2000'] * 12 + df['MONTH']
10    df['YEAR_SQUARED'] = df['YEAR_SINCE_2000'] ** 2
11
12    # One-hot encoding
13    df = pd.get_dummies(df, columns=['MONTH'], prefix='MONTH')
```

# Convert month to season, encode it, and merge with climate features

```
1 # Transform Model 1 output + extra info into Model 2 input
2 def convert_to_model2_input(model1_output, month, model2_columns):
3     max_temp, min_temp, wind_speed, temp_range, precipitation = model1_output
4
5     season = convert_month_to_season(month)
6     season_cols = ['SEASON_Spring', 'SEASON_Summer', 'SEASON_Fall', 'SEASON_Winter']
7     season_encoding = {col: 1 if col.endswith(season) else 0 for col in season_cols}
8
9     input_dict = {
10         'MAX_TEMP': max_temp,
11         'MIN_TEMP': min_temp,
12         'AVG_WIND_SPEED': wind_speed,
13         'TEMP_RANGE': temp_range,
14         'PRECIPITATION': precipitation,
15         **season_encoding
16     }
17     input_df = pd.DataFrame([input_dict])
18     input_df = input_df.reindex(columns=model2_columns, fill_value=0)
19     return input_df
```

SEASONAL  
DATA

MERGE !

ORDER !

# Define input columns for each model and run the models

```
1 def run_fire_risk_prediction(model1, model2, model1_columns, model2_columns):
2     year, month = get_user_input()
3
4     X_model1 = convert_to_model1_input(year, month, model1_columns) Generate model input
5     climate_pred = predict_climate(model1, X_model1)
6
7     X_model2 = convert_to_model2_input(climate_pred, month, model2_columns) Obtain predictions
8
9     fire_prob = predict_fire_probability(model2, X_model2) Construct Model 2 input
10    print(f"\n\n
```

Definition of required  
input columns

## RUN !!



```
1 run_fire_risk_prediction(model1, model2, model1_columns, model2_columns)
```




예측할 연도를 입력하세요 (예: 2026): 2049

예측할 월을 입력하세요 (1~12): 4

✓ 예측 결과: 2049년 4월의 산불 발생 확률은 \*\*14.22%\*\* 입니다.

# Comparison of wildfire occurrence probabilities





# 07 Significance of the project

# A Prevention-Oriented Model





Our model focuses on prevention through proactive wildfire prediction, rather than reactive responses. In this sense, it holds significant value as a strategic tool to minimize ecosystem-wide damage.


# Long-term wildfire risk trends due to climate change



# Significance of the project



Maximize the sensitivity of risk detection  
Ultimately, this project goes beyond a simple classification model – it aims to predict complex risks linked to climate change and serve as a purposeful analytical tool that can contribute to real-world wildfire response strategies.



# THANKS!

**CREDITS:** This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)